

Office hours: Mon. 1-2:30 TSRB 228

- Data structures
- Computability theory NP-hard
- 40% 4 programming (Python 2.7) projects
- 60% 2 tests
- Ungraded exercises

Solve NP-hard problems — algorithms

### **Turing test**

- Judge connected to a separate person pretending to be a machine, and a machine
- Can ask questions of the two entities, and if you can't guess which is human/machine then the machine is intelligent

**Strong AI** - emulate human behavior (beat the Turing test), broad AI (does many tasks)

**Weak/Narrow AI** - solve specific problems humans wanted to solve  
human-level, superintelligence

**Machine learning** - automated discovery of patterns in data + make decisions based on it  
- big data/big iron

**Neural networks** - started improving recently

Cognitive Revolution - instead of taking input and producing behavior, computers might want to think about the inputs

Neural Net 1943 - mimic human brain

well-defined problems: know all possible moves (chess/checkers, proofs with certain symbols) —> easier to solve by brute force

Knowledge-based AI: brittle b/c can't work with info that it can't diagnose (medical diagnosis)

ImageNet: computers learning to recognize images

Neural nets achieved 85% success rate —> efficient b/c can run in parallel

---

**Agent** - anything that perceives its envt through sensors and acts on its envt through effectors  
-interested in how to process perceptions into output actions

Agent function - mathematical description of agent's behavior in resp. to envt., mapping sensory perceptions to effector actions

Agent program - concrete implementation of agent fcn. (code)

~~Tabulation: list percept sequences, map to actions~~ <— inefficient for complex agents (naive approach)

House with two rooms, can be in A or B

Actions: move right, move left, kill person

Sensors: which room, people

- Evaluate actions according to goal of the agent (different resp. based on objective)

Objective function: defines what the goal is

-Maximize/minimize a value

-Take into account costs of and restrictions on actions

-exploration/exploitation tradeoff (look at more things vs. focus on few things)

-take into account unreliable sensors or effectors

**Observability** of agent's environment

-Fully observable if can sense everything in envt. without error

-Partially observable if limited sensors or randomness makes some info unsure

**Determinism**

-Deterministic envt only changes when agent acts, and exactly as desired

-Stochastic envt. can change at other times due to randomness, other effectors

-partially observable worlds are sensor-stochastic

-Action determinism vs. sensor determinism

**Static vs. Dynamic**

-Static world doesn't change when agent is thinking

-Dynamic world changes while you think

**Discreteness**

-World broken up into discrete pieces vs. continuous/infinite gradation of values

-Discrete vs. continuous actions, perceptions, time

**Episodic:** history doesn't matter vs. **sequential:** history matters (take future into account)

-Past events don't affect current decision

**Single-agent** vs. **multi-agent** envt. (cooperative or competitive)

## Goal-Based Agents (Search)

- Solving sequential problems (affect future decisions)
- State: unique config of relevant facts of the envt. from sensors
  - Goal state is desired state for agent —> want to figure out how to reach goal and then do it
- \*\*Need representable state (init./goal), set of acceptable states, function that accepts states

Computer sees graph with costs on different paths —> doesn't have pathfinding intuition

Use search when you don't have a better algorithm (memory-inefficient, etc.)

Picking an algorithm: uninformed/informed search, adversarial search, constraint satisfaction, conformant, Markov decision processes

**Search problem**: find a sequence of actions that transforms envt. from init. state to a goal state

state space - collection of possible states, linked to each other

Successor function: Creates new states from old, given set of actions

- Successors (single state  $s$ , set of actions  $a$ ) —>  $s'$  (new states)
- Different algorithms pick successors in different ways

—>

- Completeness: Always finds a solution if one exists **OR** can visit all states in state space
- Time complexity: Big-O, # of states created by successor fcn., worst case/average case
- Space complexity: # of states stored in memory at one timer
- Optimality: finds lowest cost/shortest solution

successor(state, actions)

```
new_states = empty set
foreach a in actions
    if a can be done to state
        add a(s) to new_states
```

1. Pick state
2. Generate successors
3. Repeat until goal
4. Remember which actions led to which state —> execute

**Issues**: inefficient (extra states visited), infinite search, not optimal, incomplete

**Fixes**: remember where you've been, be systematic

## Generic Search Algorithm

operators = [ ... list of actions ... ]

closed = nil (where you've been)

open = initial state (states we know about but haven't reached)

current = initial state

while (current isn't the goal and open != null)

    closed = closed + current

\*\*\*\*\* open = open - current union (successors(current, ops) - closed) add succ., remove closed

    current = first(open)

## Don't repeat elements in open list!

if current is the goal, success, else failure

Append to the end → open list is a queue, breadth-first search

Append to front → open list is a stack, depth-first search

Open list is priority queue → A\*, uniform cost search

**Breadth-first:** Successor fcn. gives state in alphanumeric order

-Solutions close to initial state

-Goal found → move back up decision tree for path (remember parents)

**Depth-first:**

-Follows a path as long as possible without looping

-not optimal b/c doesn't account for multiple paths to goal

Smaller closed list → more efficient process

	BFS	DFS
Complete	Complete	Yes if finite, no if infinite
Time Complexity	Number of nodes expanded (successive function calls) → branching factor ^ depth (branching factor = avg. # successors)	branching factor ^ m, m is maximum depth search will go to (generally m > depth)
Space	Keep all nodes	If state space is tree shaped, keep only ancestors (otherwise keep everything)
Optimality	yes (shortest # moves)	no

**Tradeoff: Optimality vs. space/time complexity**

### **Action Cost**

Uniform cost search (UCS) - instead of least actions, lowest total cost —> **open list is a PQ**

-g(a) - cost of moving from initial state to current state using shortest known path (minimize this)

Explore all g values starting with first element in PQ, then moves onto later ones

Note: override a node's value if it has a lower one as a child of a different parent

This is **DIJKSTRA'S ALGORITHM** —> **Optimal** and **Complete**

### **Informed Search**

Don't want to spend time looking at nodes close to init that are not intuitively efficient (ex. moving left first to get to a place on the right)

Heuristic function h() - gives an estimate of how far S is away from the goal

- Should operate in  $O(1)$ ,  $O(n)$ ,  $O(n^k)$

- Ex: Euclidean distance on map

### **Greedy Best First Search**

Sort open list (PQ) using h(s) (smaller h(s) is better) instead of g(s)

- Estimate instead of actual costs

no admissibility required

Remove item from open list, add successors with h(s) values to priority queue

Complete: yes

Time/space: same as before, but good h() will make it better

Optimality: not optimal

### **\*\*\*Best First Search\*\*\***

Merge UCS and greedy best first

Sort open list PQ on  $f(s) = g(s) + h(s)$  instead of just h(s)

Complete: yes

Time/space: exponential, but good h() helps

Optimal: depends on h()

Shortcut!

- Resort open list + revisit all descendants

- Shortcut created by "bad" heuristic

## A\*

- Best first search w/heuristic guaranteed never to create shortcuts
- Prove that  $h()$  is admissible (never overestimates distance to goal)

$f(s) = g(s) + h(s)$  but now  $h()$  must be admissible

\*\*\*Not guaranteed to find optimal solution! ( $h$  values might be too big)

Let  $h^*(s)$  be the **true cost** of state  $s$  to the goal

Also  $f^*(s) = g^*(s) + h^*(s)$  (perfect heuristic)

Heuristic is **admissible** if  $h(s) \leq h^*(s)$  for all  $s$  (heuristic never overestimates)

$h(s) = 0$  is admissible, but doesn't help at all  $\rightarrow$  UCS

**A\* is optimal (given  $h$  is admissible)  $\rightarrow$  prove admissibility**

Want admissible heuristic that is close in approximation

add start to openSet

while openSet isn't empty

    current = pop from openSet

    if current == goal

        return reconstruct\_path(current)

    closedSet.add(current)

    for every neighbor of current

        if neighbor is in closedSet, continue for loop

        gScore = current.gScore + heuristic(current, neighbor)

        if neighbor not in openSet

            openSet.add(neighbor)

        else if gScore < openSet.get(neighbor).gScore

            openSet.replace(openSet.get(neighbor), neighbor)

## Informedness

-More informed: gives value closer to  $h^*$  than other heuristics

size of total search space / avg. # states explored with heuristic  $h$  (want high informedness)

$h_a$  dominates  $h_b$  when  $h_b(n) \leq h_a(n) \leq h^*(n)$  ( $h_b$  underestimates more)

Consistency: always increasing or decreasing relative to neighbors

for all  $s_1, s_2$

$h(s_1) - h(s_2) \leq k(s_1, s_2)$

diff between heuristics less than actual cost between states

Creating a heuristic: relax problem conditions, then re-add after finding a heuristic

## Search (Randomized optimization)

heuristic, but no goal —> how to find best final state?

don't care about path

Each state is a representation, that can be searched

Naive approach: generate representation at random and test how well it does

- Issues: large hypothesis space, doesn't learn

Other approaches: Hill-climbing/greedy search, simulated annealing, genetic algos

## Hill-Climbing

y-axis is heuristic, x-axis is representations

Choose random selection from space

Initial score = heuristic(selection)

prevScore = 0

while score > prevScore

    prevScore = score

    Go through neighbors, choose highest (best) neighbor until you reach the top of the hill

If neighbors are equal, depends on neighbor function

local maxima - point with only worse neighbors

global maxima - point better than all points

Ex:

Representation: 6 bit strings

Heuristic: Shared bits with 101010

Neighbor function 1: flip pairs of adjacent bits

Neighbor function 2: flip any two bits

Use neighbor functions to find maxima with 010101, 000100, 110000

## Finding global max

Hill climb with random restarts (not guaranteed/efficient)

Works poorly in hilly or skewed sets

## Simulated Annealing

Sometimes take random steps at beginning of search, less as you continue

(See slides for pseudocode)

Temperature = 0 —> hillwalk, infinity —> random walk

-Choose a **good schedule** to make algorithm work well

## Genetic Algorithms

Instead of walking from one pt., go from many points across search space —> make better jumps

Initialization —> random mutation —> informed jumps to find global max

Mutation —> reproduce (crossover) —> evolve

mutate - replace member with random neighbor

crossover - mix two members based on heuristic value

population  $\geq 2n$

reduce - take the  $n$  best items after crossover

Grid World example

population size: 10 (arbitrary)

heuristic: Manhattan distance from goal - cost

mutation: change an action in the sequence

crossover: first half of one plan, another half of second plan

reduce: remove worst plans

## TEST: print notes

-short answer based on definitions

-worked problems (search, other algorithms)

## optimization search

need to know end (maxima) + don't care where we start or how we get there

evaluation fcn. - how good a state is (in absolute terms)

Random hill climbing - pick nearby state that is higher up until you reach a peak

Random hill climbing w/restarts - take max of multiple runs

## Simulated Annealing

always go up, unless a successor less than you appears —> go down with prob. proportional to how much less it is

-temperature - how willing you are to make a risky move (move to a lower state)

get less jumpy as algorithm continues

- slow temperature decrease —> **guaranteed to find goal**

Algorithm: generate successor from neighborhood

if  $e(\text{succ}) > e(\text{current})$  —> current becomes successor

if  $e(\text{succ}) < e(\text{current})$  then maybe current = successor

Reduce temperature

Repeat



## Genetic Algorithms

parallel hill climbing with sometimes information sharing

mutation - generate local successors

crossover - swapping information and make big jumps

Search problems w/Action stochasticity (**uncertainty** about outcomes of actions)

Episodic: solve with utility theory

Utility theory: given action  $a$  with results  $result_i(a)$  for  $i = 1$  to  $N$

Expected utility( $a$ ) = sum from  $i = 1$  to  $N$  of  $P(result_i(a)) * U(result_i(a))$

multiply probability \* utility (how likely result is \* how good it is)

**\*\*\* Utility of state is not clear in sequential problems**

Imperfect actions (ex. gridworld):

80% chance of doing what you want, 10% of drifting left, 10% of drifting right

If we fail: replan (but this is expensive!)

**Policy:**  $P: s \rightarrow a$  is a function that maps all possible states to best action from that state

**Markov Decision Process:**

1-Markov assumption: to compute best option, only need 1 piece of historical info

Creating a policy:

$S$ : set of states - sink states are those you enter + can't leave (including goals)

$s_0$ : initial state

**Transition fcn.**  $T(s, a, s')$ : probability of going from  $s$  to  $s'$  if action  $a$  is taken

$\rightarrow$  create transition table for each action

**Reward fcn.**  $R(s)$  - produces number from state (how good that state is)

-Sparse (don't get rewarded often)

-Defines optimality of behavior

Maximize reward by trying to move to state with higher reward if possible

Give some states high reward, and make them sink states  $\rightarrow$  goes to those states

Optimal policy - gets the agent highest cumulative reward

$Pi^*(s) = \text{argmax}(\text{sum of all states } T(s, a, s') * U(s'))$

- find argument (action) that gives maximum sum of transition \* utility of successor
- don't know utility, but do know reward

Utility: places might have the same score (0), but being closer to goal is obviously better

Want proximity to future reward

sink states -  $R(s) = U(s)$

**Additive utility** -  $U(s_0) = U([s_0, s_1, s_2 \dots s_n]) = R(s_0) + R(s_1) + \dots R(s_n)$   
n states in the future

for all possible future sequences.

**Discounted utility** -  $U(s_0) = U([s_0, s_1 \dots]) = R(s_0) + \gamma * R(s_1) + \gamma^2 * R(s_2) \dots$

- gamma is discount factor ( $0 < \gamma \leq 1$ )
- Trust nearby states, discount future states

**Bellman Equation** - use discounted utility

$U(s) = R(s) + \gamma * \max_a (\sum T(s, a, s') * U(s'))$

-current reward + single-step discount \* max(all possible actions: likelihood to get there \* utility)

**-recursive!!**

-utility of successors is based on their reward + utility of their successors

Solving recursion: use value iteration

-Intuition: start with random utilities and incrementally update until they reach right answer

Updating Bellman:  $U_{i+1} = R(s) + \gamma * \max_a (\sum T(s, a, s') * U(s'))$

- where  $U_i$  is guess of utility for state s after i iterations
- keep repeating function until right U is converged on